

GADTs

Advanced functional programming - Lecture 7

Wouter Swierstra and Alejandro Serrano



Today's lecture

Generalized algebraic data types (GADTs)



A datatype

```
data Tree a = Leaf  
            | Node (Tree a) a (Tree a)
```

This definition introduces:



A datatype

```
data Tree a = Leaf  
            | Node (Tree a) a (Tree a)
```

This definition introduces:

- ▶ a new datatype `Tree` of kind `* -> *`.



A datatype

```
data Tree a = Leaf
            | Node (Tree a) a (Tree a)
```

This definition introduces:

- ▶ a new datatype `Tree` of kind `* -> *`.
- ▶ two constructor functions

```
Leaf  :: Tree a
Node  :: Tree a -> a -> Tree a -> Tree a
```



A datatype

```
data Tree a = Leaf
            | Node (Tree a) a (Tree a)
```

This definition introduces:

- ▶ a new datatype `Tree` of kind `* -> *`.
- ▶ two constructor functions

```
Leaf  :: Tree a
Node  :: Tree a -> a -> Tree a -> Tree a
```

- ▶ the possibility to use the constructors `Leaf` and `Node` in patterns.



Alternative syntax

Observation

The types of the constructor functions contain sufficient information to describe the datatype.

```
data Tree :: * -> * where
  Leaf  :: Tree a
  Node  :: Tree a -> a -> Tree a -> Tree a
```

Question

What are the *restrictions* regarding the types of the constructors?



Algebraic datatypes

Constructors of an algebraic datatype T must:

- ▶ target the type T ,
- ▶ result in a simple type of kind $*$, i.e., $T\ a_1\ \dots\ a_n$ where a_1, \dots, a_n are distinct type variables.



Another example

```
data Either :: * -> * -> * where
  Left  :: a -> Either a b
  Right :: b -> Either a b
```

Both constructors produce values of type `Either a b`.

Does it make sense to lift these restrictions?



Excursion: Expression language

Imagine we're implementing a small programming language in Haskell:

```
data Expr = LitI    Int
          | LitB    Bool
          | IsZero  Expr
          | Plus    Expr Expr
          | If      Expr Expr Expr
```



Excursion: Expression language

Alternatively, we could redefine the data type as follows:

```
data Expr :: * where
  LitI    :: Int -> Expr
  LitB    :: Bool -> Expr
  IsZero  :: Expr -> Expr
  Plus    :: Expr -> Expr -> Expr
  If      :: Expr -> Expr -> Expr -> Expr
```



Syntax: concrete vs abstract

Imagined concrete syntax:

```
if isZero (0 + 1) then False else True
```

Abstract syntax:

```
If (IsZero (Plus (LitI 0) (LitI 1)))  
  (LitB False)  
  (LitB True)
```



Type errors

It is all too easy to write ill-typed expressions such as:

```
If (LitI 0) (LitB False) (LitI 1)
```

How can we prevent programmers from writing such terms?



Phantom types

At the moment, *all* expressions have the same type:

```
data Expr = LitI    Int
          | LitB    Bool
          | ...
```

We would like to distinguish between expressions of *different* types.



Phantom types

At the moment, *all* expressions have the same type:

```
data Expr = LitI    Int
          | LitB    Bool
          | ...
```

We would like to distinguish between expressions of *different* types.

To do so, we add an additional *type parameter* to our expression data type.



Phantom types

```
data Expr a = LitI    Int
            | LitB    Bool
            | IsZero   (Expr Int)
            | Plus     (Expr Int) (Expr Int)
            | If       (Expr Bool) (Expr a) (Expr a)
```

Note: the type variable *a* is never actually used in the data type for expressions.

We call such type variables *phantom types*.



Constructing well typed terms

Rather than expose the constructors of our expression language, we can instead provide a *well-typed API* for users to write terms:

```
litI :: Int -> Expr Int  
litI = LitI
```

```
plus :: Expr Int -> Expr Int -> Expr Int  
plus = Plus
```

```
isZero :: Expr Int -> Expr Bool  
isZero = IsZero
```

This guarantees that users will only ever construct well-typed terms! But what about writing an interpreter...



Evaluation

Before we write an interpreter, we need to choose the type that it returns.

Our expressions may evaluate to booleans or integers:

```
data Val = VInt    Int
         | VBool   Bool
```

Defining an interpreter now boils down to defining a function:

```
eval :: Expr a -> Val
```



Evaluation

```
eval :: Expr a -> Val
eval (LitI n)      = VInt n
eval (LitB b)      = VBool b
eval (IsZero e)    =
  case eval e of
    VInt n -> VBool (n == 0)
    _      -> error "type error"
eval (Plus e1 e2) =
  case (eval e1, eval e2) of
    (VInt n1, VInt n2) -> VInt (n1 + n2)
    _                  -> error "type error"
```



Evaluation (contd.)

- ▶ Evaluation code is mixed with code for handling type errors.
- ▶ The evaluator uses *tags* (i.e., constructors) to distinguish values – these tags are maintained and checked at run time.
- ▶ Run-time type errors can, of course, be prevented by writing a type checker or using phantom types.
- ▶ Even if we know that we only have type-correct terms, the Haskell compiler does not enforce this.



Beyond phantom types

What if we encode the type of the term in the Haskell type?

```
data Expr :: * -> * where
  LitI    :: Int -> Expr Int
  LitB    :: Bool -> Expr Bool
  IsZero  :: Expr Int -> Expr Bool
  Plus    :: Expr Int -> Expr Int -> Expr Int
  If      :: Expr Bool -> Expr a -> Expr a -> Expr a
```

Each expression has an additional *type argument*, representing the type of values it stores.



GADTs lift the restriction that all constructors must produce values of the same type.

- ▶ Constructors can have more specific return types.
- ▶ Interesting consequences for pattern matching:
 - ▶ when case-analyzing an `Expr Int`, it could not be constructed by `Bool` or `IsZero`;
 - ▶ when case-analyzing an `Expr Bool`, it could not be constructed by `Int` or `Plus`;
 - ▶ when case-analyzing an `Expr a`, once we encounter the constructor `IsZero` in a pattern, we know that we must be dealing with an `Expr Bool`;
 - ▶ ...



Evaluation revisited

```
eval :: Expr a -> a
eval (LitI n)      = n
eval (LitB b)      = b
eval (IsZero e)    = (eval e) == 0
eval (Plus e1 e2)  = eval e1 + eval e2
eval (If e1 e2 e3)
  | eval e1         = eval e2
  | otherwise       = eval e3
```

- ▶ No possibility for run-time failure (modulo undefined); no *tags* required on our values.
- ▶ Pattern matching on a GADT requires a type signature. Why?



Limitation: type signatures are required

```
data X :: * -> * where  
  C  :: Int -> X Int  
  D  :: X a  
  E  :: Bool -> X Bool
```

```
f (C n) = [n]           -- (1)
```

```
f D     = []            -- (2)
```

```
f (E n) = [n]           -- (3)
```



Limitation: type signatures are required

```
data X :: * -> * where
  C  :: Int -> X Int
  D  :: X a
  E  :: Bool -> X Bool
```

```
f (C n) = [n]           -- (1)
```

```
f D     = []           -- (2)
```

```
f (E n) = [n]           -- (3)
```

What is the type of `f`, with/without (3)? What is the (probable) desired type?

```
f :: X a -> [Int]       -- (1) only
```

```
f :: X b -> [c]         -- (2) only
```

```
f :: X a -> [Int]       -- (1) + (2)
```



Extending our language

Let us extend the expression types with pair construction and projection:

```
data Expr :: * -> * where
```

```
...
```

```
Pair      :: Expr a -> Expr b -> Expr (a, b)
```

```
Fst       :: Expr (a,b)      -> Expr a
```

```
Snd       :: Expr (a,b)      -> Expr b
```

For Fst and Snd, the type of the non-projected component is 'hidden' – that is, it is not visible from the type of the compound expression.



Evaluation again

```
eval :: Expr a -> a  
eval ...
```

```
eval (Pair x y) = (eval x, eval y)  
eval (Fst p)    = fst (eval p)  
eval (Snd p)    = snd (eval p)
```



GADTs have become one of the more popular Haskell extensions.

The ‘classic’ example for motivating GADTs is interpreters for expression languages, such as the one we have seen here.

However, these richer data types offer many other applications.

In particular, they let us *program* with types in interesting new ways.



Prelude.head: empty list

```
> myComplicatedFunction 42 "inputFile.csv"  
*** Exception: Prelude.head: empty list
```

Can we use the *type system* to rule out such exceptions before a program is run?



Prelude.head: empty list

```
> myComplicatedFunction 42 "inputFile.csv"  
*** Exception: Prelude.head: empty list
```

Can we use the *type system* to rule out such exceptions before a program is run?

To do so, we'll introduce a new list-like datatype that records the *length* of the list in its *type*.



Natural numbers and vectors

Natural numbers can be encoded as types – no constructors are required.

```
data Zero
data Succ a
```

Vectors are lists with a fixed number of elements:

```
data Vec :: * -> * -> * where
  Nil    ::                               Vec a Zero
  Cons   :: a -> Vec a n -> Vec a (Succ n)
```



Type-safe head and tail

```
head :: Vec a (Succ n) -> a  
head (Cons x xs) = x
```

```
tail :: Vec a (Succ n) -> Vec a n  
tail (Cons x xs) = xs
```

Question

Why is there no case for Nil is required?



Type-safe head and tail

```
head :: Vec a (Succ n) -> a  
head (Cons x xs) = x
```

```
tail :: Vec a (Succ n) -> Vec a n  
tail (Cons x xs) = xs
```

Question

Why is there no case for Nil is required?

Actually, a case for Nil results in a type error.



More functions on vectors

```
map :: (a -> b) -> Vec a n -> Vec b n
map f Nil          = Nil
map f (Cons x xs)  = Cons (f x) (map f xs)
```

```
zipWith :: (a -> b -> c) ->
          Vec a n -> Vec b n -> Vec c n
zipWith op Nil      Nil      = Nil
zipWith op (Cons x xs) (Cons y ys) =
  Cons (op x y) (zipWith op xs ys)
```

We can require that the two vectors have the same length!

This lets us rule out bogus cases.



Yet more functions on vectors

```
snoc :: Vec a n -> a -> Vec a (Succ n)
snoc Nil          y = Cons y Nil
snoc (Cons x xs)  y = Cons x (snoc xs y)
```

```
reverse :: Vec a n -> Vec a n
reverse Nil          = Nil
reverse (Cons x xs)  = snoc xs x
```

What about appending two vectors, analogous to the (++) operation on lists?



Problematic functions

- ▶ What is the type of our append function?

`vappend :: Vec a m -> Vec a n -> Vec a ???`



Problematic functions

- ▶ What is the type of our append function?

`vappend :: Vec a m -> Vec a n -> Vec a ???`

How can we add two *types*, n and m ?



Problematic functions

- ▶ What is the type of our append function?

```
vappend :: Vec a m -> Vec a n -> Vec a ???
```

How can we add two *types*, n and m ?

- ▶ Suppose we want to convert from lists to vectors:

```
fromList :: [a] -> Vec a n
```

Where does the type variable n come from? What possible values can it have?



Writing vector append

There are multiple options to solve that problem:

- ▶ construct explicit evidence,
- ▶ use a type family (more on that in the next lecture).



Explicit evidence

Given two 'types' n and m , what is their sum?

We can define a GADT describing the *graph* of addition:

```
data Sum :: * -> * -> * -> * where
  SumZero :: Sum Zero n n
  SumSucc  :: Sum m n s -> Sum (Succ m) n (Succ s)
```



Explicit evidence

Given two 'types' n and m , what is their sum?

We can define a GADT describing the *graph* of addition:

```
data Sum :: * -> * -> * -> * where
  SumZero :: Sum Zero n n
  SumSucc  :: Sum m n s -> Sum (Succ m) n (Succ s)
```

Using this function, we can now define append as follows:

```
append :: Sum m n s
        -> Vec a m -> Vec a n -> Vec a s
append SumZero Nil ys = ys
append (SumSucc p) (Cons x xs) ys =
  Cons x (append p xs ys)
```



Passing explicit evidence

This approach has one major disadvantage: we must construct the evidence, the values of type $\text{Sum } n \ m \ p$, by hand every time we wish to call `append`.

We could use a multi-parameter type class with functional dependencies to automate this construction...



Converting between lists and vectors

It is easy enough to convert from a vector to a list:

```
toList :: Vec a n -> [a]
toList Nil           = []
toList (Cons x xs)  = x : toList xs
```

This simply discards the type information we have carefully constructed.



Converting between lists and vectors

It is easy enough to convert from a vector to a list:

```
toList :: Vec a n -> [a]
toList Nil           = []
toList (Cons x xs)  = x : toList xs
```

This simply discards the type information we have carefully constructed.



Converting between lists and vectors

Converting in the other direction, however is not as easy:

```
fromList :: [a] -> Vec a n
fromList []      = Nil
fromList (x:xs) = Cons x (fromList xs)
```

Question

Why doesn't this definition type check?



Converting between lists and vectors

Converting in the other direction, however is not as easy:

```
fromList :: [a] -> Vec a n
fromList []      = Nil
fromList (x:xs) = Cons x (fromList xs)
```

Question

Why doesn't this definition type check?

The type says that the result must be polymorphic in n , that is, it returns a vector of *any* length, rather than a vector of a specific (unknown) length.



From lists to vectors

We can

- ▶ specify the length of the vector being constructed in a separate argument,
- ▶ hide the length using an *existential* type.



From lists to vectors (contd.)

Suppose we simply pass in a regular natural number, `Nat`:

```
fromList :: Nat -> [a] -> Vec a n
fromList Zero      []      = Nil
fromList (Succ n) (x:xs) = Cons x (fromList n xs)
fromList _         _       = error "wrong length!"
```



From lists to vectors (contd.)

Suppose we simply pass in a regular natural number, `Nat`:

```
fromList :: Nat -> [a] -> Vec a n
fromList Zero      []      = Nil
fromList (Succ n)  (x:xs) = Cons x (fromList n xs)
fromList _         _       = error "wrong length!"
```

This still does not solve our problem – there is no connection between the natural number that we are passing and the `n` in the return type.



Singletons

We need to reflect type-level natural numbers on the value level.

To do so, we define a (yet another) variation on natural numbers:

```
data SNat :: * -> * where
  SZero :: SNat Zero
  SSucc  :: SNat n -> SNat (Succ n)
```

This is a *singleton type* – for any n , the type $\text{SNat } n$ has a single inhabitant (the number n).



From lists to vectors

```
data SNat :: * -> * where
  SZero  :: SNat Zero
  SSucc  :: SNat n -> SNat (Succ n)

fromList :: SNat n -> [a] -> Vec a n
fromList SZero      []      = Nil
fromList (SSucc n) (x:xs) = Cons x (fromList n xs)
fromList _          _      = error "wrong length!"
```

Question

This function may still fail dynamically. Why?



From lists to vectors

We can

- ▶ specify the length of the vector being constructed in a separate argument,
- ▶ hide the length using an *existential* type.

What about the second alternative?



From lists to vectors

We can define a wrapper around vectors, *hiding* their length:

```
data VecAnyLen :: * -> * where  
  VecAnyLen :: Vec a n -> VecAnyLen a
```

A value of type `VecAnyLen a` stores a vector of *some* length with values of type `a`.



From lists to vectors

We can convert any list to a vector of some length as follows:

```
fromList :: [a] -> VecAnyLen a
fromList []      = VecAnyLen Nil
fromList (x:xs) =
  case fromList xs of
    VecAnyLen ys -> VecAnyLen (Cons x ys)
```



From lists to vectors

We can combine the two approaches and include a `SNat` in the packed type:

```
data VecAnyLen :: * -> * where
  VecAnyLen :: SNat n -> Vec a n -> VecAnyLen a
```

Question

How does the conversion function change?



Calling functions on vectors

Given two vectors $xs : \text{Vec } a \ n$ and $ys : \text{Vec } a \ m$.

Suppose I want to zip these vectors together using:

$\text{zipVec} :: \text{Vec } a \ n \rightarrow \text{Vec } b \ n \rightarrow \text{Vec } (a,b) \ n$

Question

What happens when I call `zip xs ys`?



Comparing the length of vectors

We can define a boolean function that checks when two vectors have the same length

```
equalLength :: Vec a m -> Vec b n -> Bool
equalLength Nil Nil = True
equalLength (Cons _ xs) (Cons _ ys) =
    equalLength xs ys
```



Comparing the length of vectors

Such a function is not very useful...

Suppose I want to use this to check the lengths of my vectors:

```
if equalLength xs ys
  then zip xs ys
  else error "Wrong lengths"
```

Question

Will this type check?



Comparing the length of vectors

Such a function is not very useful...

Suppose I want to use this to check the lengths of my vectors:

```
if equalLength xs ys
  then zip xs ys
  else error "Wrong lengths"
```

Question

Will this type check?

No! Just because `equalLength xs ys` returns `True`, does not guarantee that `m` and `n` are equal...

How can we enforce that two types are indeed equal?



Equality type

Just as we saw for the Sum type, we can introduce a GADT that represents a 'proof' that two types are equal:

```
data Equal :: * -> * -> * where  
  Refl :: Equal a a
```



Equality type

Just as we saw for the Sum type, we can introduce a GADT that represents a 'proof' that two types are equal:

```
data Equal :: * -> * -> * where  
  Refl :: Equal a a
```

We can even 'prove' properties of our equality relation:

```
refl  :: Equal a a  
sym   :: Equal a b -> Equal b a  
trans :: Equal a b -> Equal b c -> Equal a c
```



Equality type

Just as we saw for the Sum type, we can introduce a GADT that represents a 'proof' that two types are equal:

```
data Equal :: * -> * -> * where  
  Refl :: Equal a a
```

We can even 'prove' properties of our equality relation:

```
refl  :: Equal a a  
sym   :: Equal a b -> Equal b a  
trans :: Equal a b -> Equal b c -> Equal a c
```

Question

How are these functions defined? What happens if you don't pattern match on the Refl constructor?



Equality type

Instead of returning a boolean, we can now provide evidence that the length of two vectors is equal:

```
eqLength :: Vec a m -> Vec b n -> Maybe (Equal m n)
eqLength Nil Nil = Just Refl
eqLength (Cons x xs) (Cons y ys) =
  case eqLength xs ys of
    Just Refl -> Just Refl
    Nothing -> Nothing
eqLength _ _ = Nothing
```



Using equality

```
test :: Vec a m -> Vec b (Succ n) -> Maybe (a,b)
test xs ys =
  case eqLength xs ys
    of Just Refl -> head (zip xs ys)
       _         -> Nothing
```

Question

Why does this type check?



Expressive power of equality

The equality type can be used to encode other GADTs.

Recall our expression example using phantom types:

```
data Expr a =  
  LitI    Int  
| LitB    Bool  
| IsZero  (Expr a)  
| Plus    (Expr a) (Expr a)  
| If      (Expr a) (Expr a) (Expr a)
```



Expressive power of equality

We can replace this with a phantom type

```
data Expr a =  
    LitI    (Equal a Int) Int  
  | LitB    (Equal a Bool) Bool  
  | IsZero  (Equal a Bool) (Equal b Int) (Expr b)  
  | Plus    (Equal a Int) (Expr a) (Expr a) (Expr a)  
  ...
```



Safe vs unsafe coercions

Using our equality function we can safely coerce between types:

```
coerce :: Equal a b -> a -> b  
coerce Refl x = x
```

Question

Why does this type check?



Safe vs unsafe coercions

Using our equality function we can safely coerce between types:

```
coerce :: Equal a b -> a -> b  
coerce Refl x = x
```

Question

Why does this type check?

Question

What about this definition:

```
coerce :: Equal a b -> a -> b  
coerce p x = x
```



Aside: irrefutable patterns

Haskell also allows irrefutable patterns:

`lazyHead ~(x:xs) = x`

This does not force the list to weak head normal form.



Aside: irrefutable patterns

In tandem with GADTs this is particularly dangerous:

```
coerceL :: Equal a b -> a -> b  
coerceL ~Refl x = x
```

Question

How could this cause well-typed program to crash with a type error?



Aside: irrefutable patterns

In tandem with GADTs this is particularly dangerous:

```
coerceL :: Equal a b -> a -> b  
coerceL ~Refl x = x
```

Question

How could this cause well-typed program to crash with a type error?

```
foo :: Bool -> Int  
foo b = coerceL undefined b
```

Apparently unrelated language features may interact in unexpected ways!



Outlook generic programming: Reflecting types

We can even use GADTs to *reflect* types themselves as data:

```
data Type :: * -> * where
  INT  :: Type Int
  BOOL :: Type Bool
  LIST :: Type a -> Type [a]
  PAIR :: Type a -> Type b -> Type (a,b)
```



Safe dynamically typed values

We can define dynamically typed values by packing up a type representation with a value:

```
data Dynamic :: * where  
  Dyn :: Type a -> a -> Dynamic
```



Safe dynamically typed values

We can define dynamically typed values by packing up a type representation with a value:

```
data Dynamic :: * where
  Dyn :: Type a -> a -> Dynamic
```

To unwrap these values safely, we check whether the types line up as expected:

```
coerce :: Type a -> Dynamic -> Maybe a
coerce t (Dyn t' x) =
  case eqType t t'
    of Just Refl -> Just x
       _         -> Nothing
```



Generic programming

We can also define new functions *by induction on the type structure*:

$f :: \text{Type } a \rightarrow \dots a \dots$

In this way, we can define our own versions of functions such as `show`, `read`, `equality`, etc.



Summary

- ▶ GADTs can be used to encode advanced properties of types in the type language.
- ▶ We end up mirroring expression-level concepts on the type level (e.g. natural numbers).
- ▶ GADTs can also represent data that is computationally irrelevant and just guides the type checker (equality proofs, evidence for addition).
Such information could ideally be erased, but in Haskell, we can always cheat via undefined `:: Equal Int Bool...`

